


S1b: with this memoryless reviewer, excluding caught cheaters (not using audit results) produced the safety



Misreporters under-report by half ($d = 0.5$); 64 contexts per game. 'Audit used' = the audited agent's true request enters that step's decision only (no memory, no sanction). 'Exclude' = S1 rule (6 steps out + fine), which also removes extraction, so it can go below the honest level (dotted line = everyone honest / full verification). Note different y-scales by panel.